

Module 6: Deadlock

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

Operating Systems

6.1

Lafuze 2012 - Silberschatz and Galvin

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- Example: System has 2 tape drives. P_1 and P_2 each hold one tape drive and each needs another one.
Solution with 2 semaphores A and B , initialized to 1

P_0
 $wait(A);$
 $wait(B);$

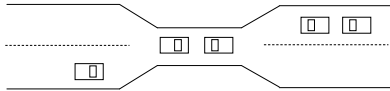
P_1
 $wait(B)$
 $wait(A)$

Operating Systems

6.2

Lafuze 2012 - Silberschatz and Galvin

Bridge Crossing Example




- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.

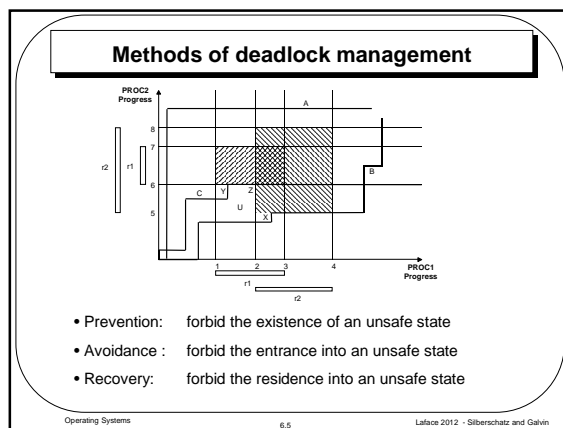
Operating Systems

6.3

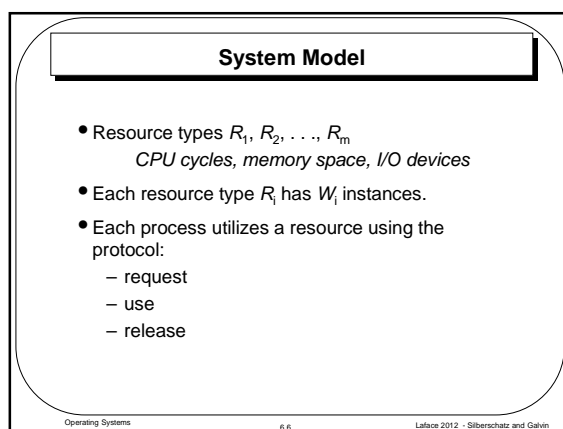
Lafuze 2012 - Silberschatz and Galvin

1





- Prevention: forbid the existence of an unsafe state
- Avoidance : forbid the entrance into an unsafe state
- Recovery: forbid the residence into an unsafe state



Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Operating Systems

6.7

Lafuze 2012 - Silberschatz and Galvin

Resource-Allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- request edge – directed edge $P_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow P_i$

Operating Systems

6.8

Lafuze 2012 - Silberschatz and Galvin

Resource-Allocation Graph (Cont.)

- Process



- Resource Type with 4 instances



- P_i requests instance of R_j



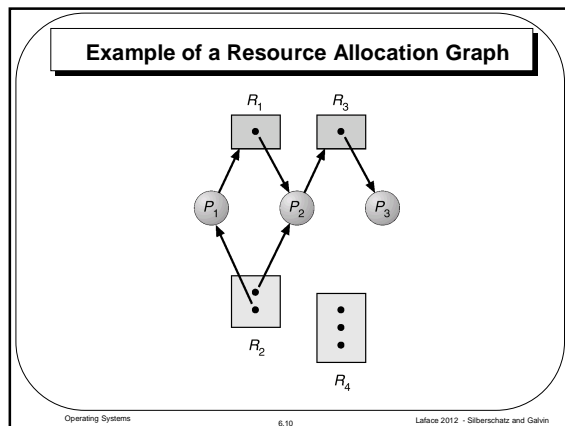
- P_i is holding an instance of R_j

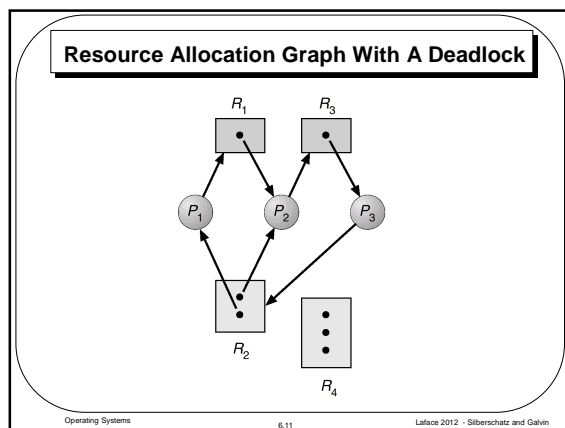



Operating Systems

6.9

Lafuze 2012 - Silberschatz and Galvin







Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.

Operating Systems

6.13

Lafuze 2012 - Silberschatz and Galvin

Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
- Allow the system to enter a deadlock state and then recover.
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX.

Operating Systems

6.14

Lafuze 2012 - Silberschatz and Galvin

Deadlock Prevention

Restrain the ways request can be made.
Violate the conditions that are necessary for deadlock.

- Mutual Exclusion – not required for sharable resources; must hold for nonsharable resources.

Operating Systems

6.15

Lafuze 2012 - Silberschatz and Galvin

Deadlock Prevention

- Hold and Wait – must guarantee that whenever a process requests a resource, it does not hold any other resources.
 - **Request All First** (RAF): Require process to request and be allocated **all** its resources before it begins execution, or allow process to request resources only when the process has none.
 - **Release Before Request** (RBR): process can request and wait for resources only if they don't keep resources already acquired. They must release their resources before they can require new resources.

Drawbacks:
Low resource utilization; starvation possible.

Operating Systems
6.16
Lafuze 2012 - Silberschatz and Galvin

Deadlock Prevention

- No Preemption
 - If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released (preemption).
 - Preempted resources are added to the list of resources for which the process is waiting.
 - Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

Operating Systems
6.17
Lafuze 2012 - Silberschatz and Galvin

Deadlock Prevention

- **Circular Wait** – *Hierarchical Resource Usage* (HRU).
- Impose a total ordering of all resource types, associating each resource to an integer number, for example:

disks = 1

tapes = 5

printers = 12
- and require that each process requests resources in an increasing order of enumeration.

Operating Systems
6.18
Lafuze 2012 - Silberschatz and Galvin

HRU

- A process can require an instance of resource r_k
 - if $F(r_k) > F(r_i)$,
where r_i is a type of resource obtained before r_k .
- Otherwise, the process must release all the resources of type r_h such that $F(r_k) \leq F(r_h)$.

Operating Systems
6.19
Lafuze 2012 - Silberschatz and Galvin

Hierarchical Resource Usage

Let's suppose that there exists a set of processes $P = \{P_0, P_1, \dots, P_n\}$ in circular wait \Rightarrow

$$F(r_k) < F(r_{k+1}), \forall k = 0 \dots n-1 \quad \text{i.e.}$$

$$F(r_0) < F(r_1) < \dots < F(r_n) < F(r_0)$$

$$F(r_0) < F(r_0)$$

this is clearly absurd.

- A condition of circular wait, thus, cannot happen.

Operating Systems
6.20
Lafuze 2012 - Silberschatz and Galvin

Deadlock Avoidance

Requires that the system has some additional *a priori* information available.

- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes.

Operating Systems
6.21
Lafuze 2012 - Silberschatz and Galvin

Banker's Algorithm: State

Cash

2
4 (4)
2 (1)
2 (7)

P

Q

R

Acquired Resources (remaining requests)

Operating Systems

6.22

Lafuze 2012 - Silberschatz and Galvin

Banker's Algorithm: Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.
- System is in safe state if there exists a safe sequence of resource allocation for all processes.
- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

Operating Systems

6.23

Lafuze 2012 - Silberschatz and Galvin

Banker's Algorithm

- Each process declare the *maximum number* of resources it will need
- When it requires a resource it can be blocked for a limited amount of time
- When it obtain the resources, it will release them in a finite amount of time

Complexity $O(m \times n^2)$, where m is the number of resources and n the number of processes.

Operating Systems

6.24

Lafuze 2012 - Silberschatz and Galvin

Example of safe state											
Cash	2	Cash	4	Cash	8	Cash	10				
P	4 (4)	P	4 (4)	P	- (-)	P	- (-)				
Q	2 (1)	Q	- (-)	Q	- (-)	Q	- (-)				
R	2 (7)	R	2 (7)	R	2 (7)	R	- (-)				

Operating Systems

6.25


Lafuze 2012 - Silberschatz and Galvin

Example of unsafe state											
Cash	2	Cash	1	Cash	3						
P	4 (4)	P	4 (4)	P	4 (4)						
Q	2 (1)	Q	2 (1)	Q	- (-)						
R	2 (7)	R	3 (6)	R	3 (6)						

Operating Systems

6.26

Lafuze 2012 - Silberschatz and Galvin



Operating Systems

6.27

Lafuze 2012 - Silberschatz and Galvin

Deadlock Detection

- Allow system to enter deadlock state
- Maintain *wait-for* graph
 - Nodes are processes.
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j .
- Periodically, or when a process requires a resource, invoke an algorithm that searches for a cycle in the graph
- If a cycle exists the a recovery scheme is applied

An algorithm to detect a cycle in a graph requires $O(n^2)$ operations, where n is the number processes.

Operating Systems

6.28

Lafuze 2012 - Silberschatz and Galvin

Recovery from Deadlock: Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer to completion.
 - Resources the process has used.
 - Resources process needs to complete.
 - How many processes will need to be terminated.
 - Is process interactive or batch?

Operating Systems

6.29

Lafuze 2012 - Silberschatz and Galvin

Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process fro that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.

Operating Systems

6.30

Lafuze 2012 - Silberschatz and Galvin

Exampe of rollback scheme

- Threads that acquire 3 locks:
 - one thread in sequence 1, 2, and 3
 - the second thread in opposite sequence: 3, 2, and 1
- Typical condition of possible deadlock

backoff.c

Operating Systems

6.31

Laface 2012 - Silberschatz and Galvin

Deadlock using IPC

```
graph LR; P1((P1)) --> B1[B1]; B1 --> P2((P2)); P2 --> B2[B2]; B2 --> P3((P3)); P3 --> B3[B3]; B3 --> P1
```

P1	P2	P3
... send (B1,m)	... send (B2,m)	... send (B3,m)
... send (B1,m)	... send (B2,m)	... send (B3,m)
... receive (B3,m)	... receive (B1,m)	... receive (B2,m)

time

m = message
Bi = i-th communication buffer

Operating Systems

6.32

Laface 2012 - Silberschatz and Galvin

Deadlock using IPC

```
graph LR; P((P)) --> PR[PR]; PR --> R((R)); R --> RS[RS]; RS --> S((S)); S --> QS[QS]; QS --> Q((Q)); Q --> PQ[PQ]; PQ --> P
```

P	Q	R	S
... send (PR,m) send (RS,m)	...
... send (PR,m)	... receive (PQ,m)	... send (RS,m)	... receive (QS,m)
... send (PQ,m)	... send (QS,m)	... receive (PR,m)	... receive (RS,m)

Operating Systems

6.33

Laface 2012 - Silberschatz and Galvin

11