

Module 4: Processes

- Process Concept
- Process Scheduling
- Operation on Processes
- Cooperating Processes
- Interprocess Communication

Operating System Concepts4s.1Lafaze 2012 -Silberschatz and Galvin ©1999

Definitions

- **Algorithm:** a logical procedure that in a finite number of steps solves a problem
- **Program:** formal expression of an algorithm by means of a programming language

A programming language needs a *run-time support*

Operating System Concepts4s.2Lafaze 2012 -Silberschatz and Galvin ©1999

Process Concept

- **Process:** a **sequence** of operations performed by a program in execution on a given set of input data.
- The temporal behavior of a process can be analyzed through its *trace*
- **Sequential process:** is a deterministic process
 - Given a set of input data always produces the same output independently from
 - the start executions time
 - the execution speed
 - the number of active processes on the same system

Operating System Concepts4s.3Lafaze 2012 -Silberschatz and Galvin ©1999

Process Concept

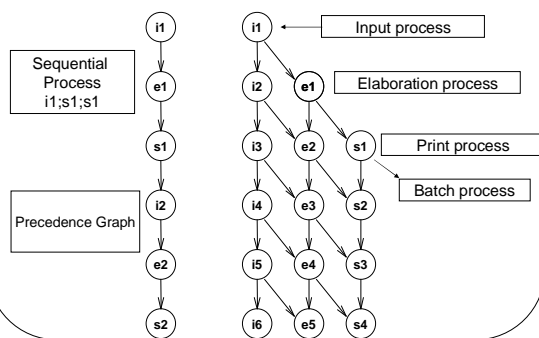
- Process – a sequence of operations performed by a program in execution on a given set of input data.
- The temporal behavior of a process can be analyzed through its *trace*
- A process includes:
 - program counter
 - stack
 - data section

Operating System Concepts

4a.4

Lafaze 2012 · Silberschatz and Galvin ©1999

Sequential and Concurrent Processes



Operating System Concepts

4a.5

Lafaze 2012 · Silberschatz and Galvin ©1999

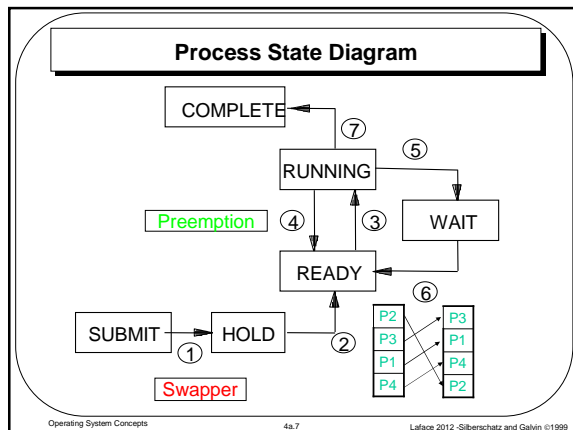
Process State

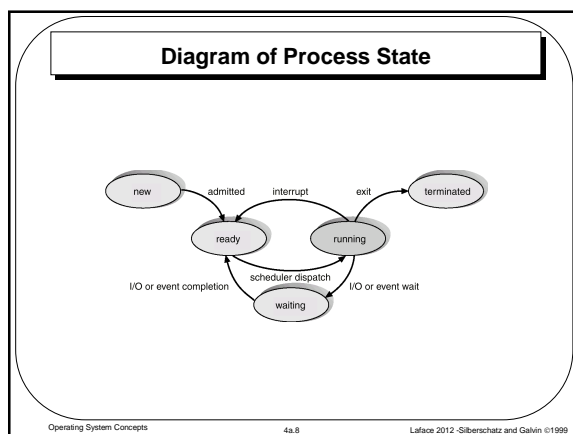
- As a process executes, it changes *state*
 - submit: The process is being created.
 - hold: The process wait for main memory resource.
 - running: Instructions are being executed.
 - waiting: The process is waiting for some event to occur.
 - ready: The process is waiting to be assigned to a processor.
 - complete: The process has finished execution.

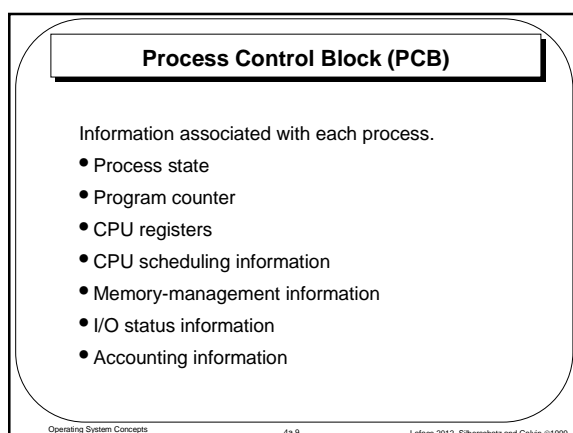
Operating System Concepts

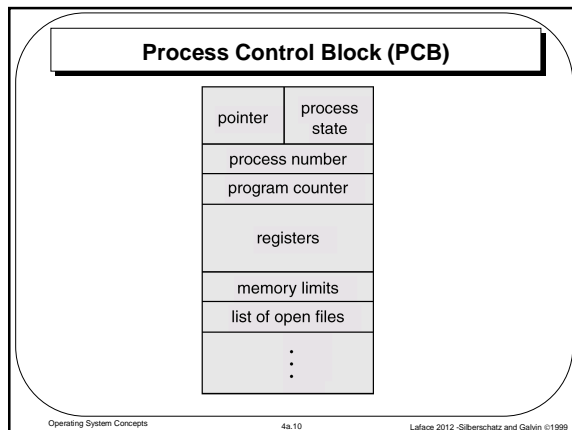
4a.6

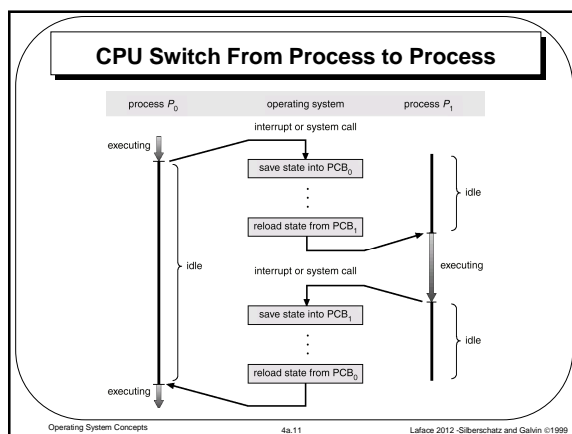
Lafaze 2012 · Silberschatz and Galvin ©1999







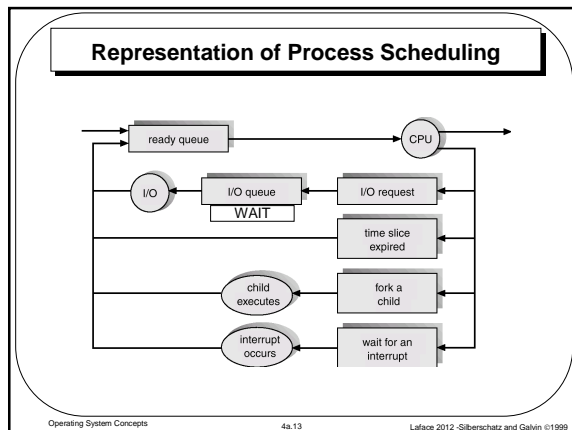




Context Switching

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Time devoted to context switching is *overhead*, i.e. work not directly useful for any process.
- Context switching time must be minimum. It is dependent on HW support.

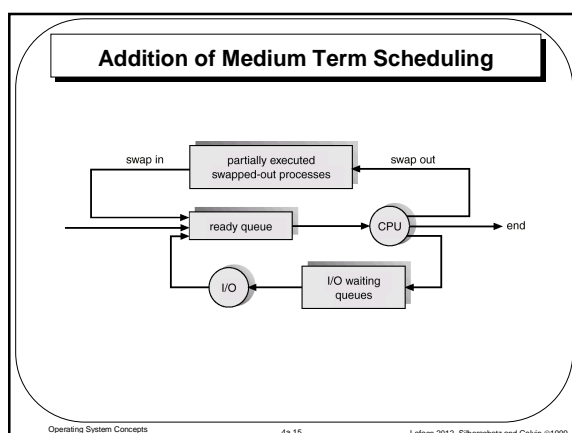
Operating System Concepts 4a.12 Lafuze 2012 -Silberschatz and Galvin ©1999



Schedulers

- Long-term scheduler (or job scheduler) – selects which processes should be brought into the ready queue.
- Short-term scheduler (or CPU scheduler) – selects which process should be executed next and allocates CPU.

Operating System Concepts 4a.14 Lafuze 2012 · Silberschatz and Galvin ©1999



Schedulers (Cont.)

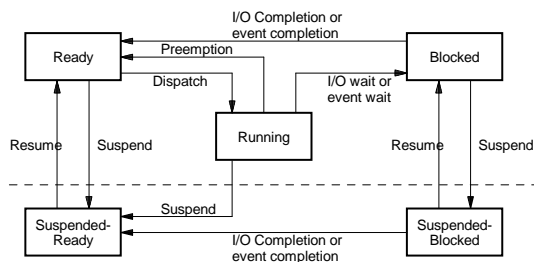
- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow).
- The medium-term scheduler controls the *degree of multiprogramming*.
- Processes can be described as either:
 - *I/O-bound process* – spends more time doing I/O than computations, many short CPU bursts.
 - *CPU-bound process* – spends more time doing computations; few very long CPU bursts.

Operating System Concepts

4s.16

Lafaze 2012 · Silberschatz and Galvin ©1999

Process State Diagram (more detailed)




Operating System Concepts

4s.17

Lafaze 2012 · Silberschatz and Galvin ©1999

Process State Diagram - UNIX



Operating System Concepts

4s.18

Lafaze 2012 · Silberschatz and Galvin ©1999

Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes.
- Resource sharing
 - Parent and children share all resources.
 - Children share subset of parent's resources.
 - Parent and child share no resources.
- Execution
 - Parent and children execute concurrently.
 - Parent waits until children terminate.

Operating System Concepts

4s.19

Lafaze 2012; Silberschatz and Galvin ©1999

Process Creation (Cont.)

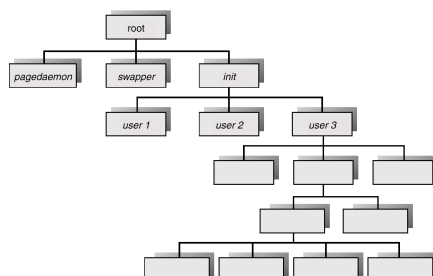
- Address space
 - Child duplicate of parent.
 - Child has a program loaded into it.
- UNIX examples
 - **fork** system call creates new process
 - **execve** system call used after a **fork** to replace the process' memory space with a new program.

Operating System Concepts

4s.20

Lafaze 2012; Silberschatz and Galvin ©1999

A Tree of Processes On a UNIX System



Operating System Concepts

4s.21

Lafaze 2012; Silberschatz and Galvin ©1999

fork

fork0.c

alarm.c Sequential Alarm Clock

alarm-fork.c Concurrent Alarm Clock

fork3.c Erroneous generation of K childrens

fork-K-children.c

Operating System Concepts
4a.22
Lafaze 2012 ·Silberschatz and Galvin ©1999

Process Address Space

fork1.c

Separate address space

Operating System Concepts
4a.23
Lafaze 2012 ·Silberschatz and Galvin ©1999

Process Termination

- Process terminates if it executes the system call **exit** or the last statement.
- The child can pass a byte to the parent (via **wait**).
- The parent receive the information by means of the **wait** argument.
- The OS kernel recovers the resources used by the child process (via **wait**).

Operating System Concepts
4a.24
Lafaze 2012 ·Silberschatz and Galvin ©1999

Fork, wait and exit

fork2.c Fork wait and exit

wait.c

Operating System Concepts4a.25Lafaze 2012 -Silberschatz and Galvin ©1999

fork - exec

execlp.c

execve.c

Operating System Concepts4a.26Lafaze 2012 -Silberschatz and Galvin ©1999

UNIX Shell Skeleton

```

while ( TRUE ){
    write_prompt;
    read_command( command, parameters );
    if ( fork() != 0 )
        /* parent process */
        wait ( &status ); /* wait */
    else
        /* child process */
        execve( command, parameters );
}

```

Operating System Concepts4a.27Lafaze 2012 -Silberschatz and Galvin ©1999

UNIX Shell Skeleton

Background command
ls -R / > /dev/null &

```

while ( TRUE ){
    write_prompt;
    read_command( command, parameters );
    if ( fork() != 0 )
        /* parent process */
        /* wait ( &status ); don't wait*/
    else
        /* child process */
        execve( command, parameters );
}

```

Operating System Concepts
4a.28
Laface 2012 -Silberschatz and Galvin ©1999

Process Address Space

Parent	Child before execve	Child after execve
x = 4; x = 10 . . wait(&status); .	x = 4; x = 10; . . execve('ls',...) .	ls command program . . . exit(0); .

Operating System Concepts
4a.29
Laface 2012 -Silberschatz and Galvin ©1999

Cooperating Processes

- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Operating System Concepts
4a.30
Laface 2012 -Silberschatz and Galvin ©1999

Producer-Consumer Problem


- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process.
 - *unbounded-buffer* places no practical limit on the size of the buffer.
 - *bounded-buffer* assumes that there is a limited buffer size.

Operating System Concepts

4a.31

Lafuze 2012 · Silberschatz and Galvin ©1999

Circular Buffer



Operating System Concepts

4a.32

Lafuze 2012 · Silberschatz and Galvin ©1999

Bounded-Buffer – Shared-Memory Solution

```
#define N 10
typedef ..... Item ; Item buffer[N];
in = 0, out=0;
```

- Producer process

```
while (true) {
```

```
...
```

```
/* produces an item in nextp */
```

```
...
```

```
while ((in+1) % N == out) ;
```

```
buffer[in] = nextp;
```

```
in = (in+1) % N;
```

```
}
```

!!!!

Operating System Concepts

4a.33

Lafuze 2012 · Silberschatz and Galvin ©1999

Bounded-Buffer (Cont.)

- Consumer Process

```

while (true) {
    while (in == out) ;
    nextc = buffer[out];
    out = (out+1) % N;
    ...
    /* consumes the item in nextc */
    ...
}

```
- The solution is correct for $N \neq 1$.
- But it wastes one buffer element.

Operating System Concepts

4a.34

Lafaze 2012 · Silberschatz and Galvin ©1999

Threads

- A *thread* (or *lightweight process*) is a basic unit of CPU utilization; it consists of:
 - program counter
 - register set
 - stack space
- A thread shares with its peer threads its:
 - code section
 - data section
 - operating-system resources
collectively known as a the **task domain**.
- A traditional or *heavyweight* process is equal to a task with one thread

Operating System Concepts

4a.35

Lafaze 2012 · Silberschatz and Galvin ©1999

Threads

Operating System Concepts

4a.36

Lafaze 2012 · Silberschatz and Galvin ©1999

Threads (Cont.)

- In a multiple threaded task, while one server thread is blocked and waiting, a second thread in the same task can run.
- Cooperation of multiple threads in same job confers higher throughput and improved performance.
- Threads allow
 - concurrent execution of parts of the same application
 - service of several clients using copies of a code section
 - sequential processes to make blocking system calls while also achieving parallelism
 - exploiting multiprocessor boards

Operating System Concepts

4a.37

Lafaze 2012 ·Silberschatz and Galvin ©1999

Multiple Threads within a Task

Operating System Concepts

4a.38

Lafaze 2012 ·Silberschatz and Galvin ©1999

Threads (Cont.)

- Kernel-supported threads (Mach and OS/2).
- User-level threads; supported above the kernel, via a set of library calls at the user level (Project Andrew from CMU).
- Hybrid approach implements both user-level and kernel-supported threads (Solaris 2).

Operating System Concepts

4a.39

Lafaze 2012 ·Silberschatz and Galvin ©1999

User threads

- Pros
 - Fast context switching between threads of the same task.
 - Can be implemented on top of any kernel
- Cons
 - If a thread performs a blocking system call, all the threads of the same task are blocked
 - A single thread per task running even on a multiprocessor system

Operating System Concepts4a.40Lafaze 2012 -Silberschatz and Galvin ©1999

Kernel threads

- Pros
 - The ready threads can be scheduled even if they belong to the task of a thread that has issued a blocking system call
 - Multiple thread per task can be executed on a multiprocessor system
- Cons
 - More expensive context switching because it needs the passage to kernel mode and kernel support
 - Constraints on the maximum number of thread per task/system

Operating System Concepts4a.41Lafaze 2012 -Silberschatz and Galvin ©1999

Hybrid approach - Solaris 2 Threads

Operating System Concepts4a.42Lafaze 2012 -Silberschatz and Galvin ©1999

Pthreads Examples

self.c

creation-and-join.c

try-exit.c

try-common-memory

Operating System Concepts

4a.43

Lafuze 2012 · Silberschatz and Galvin 11999
